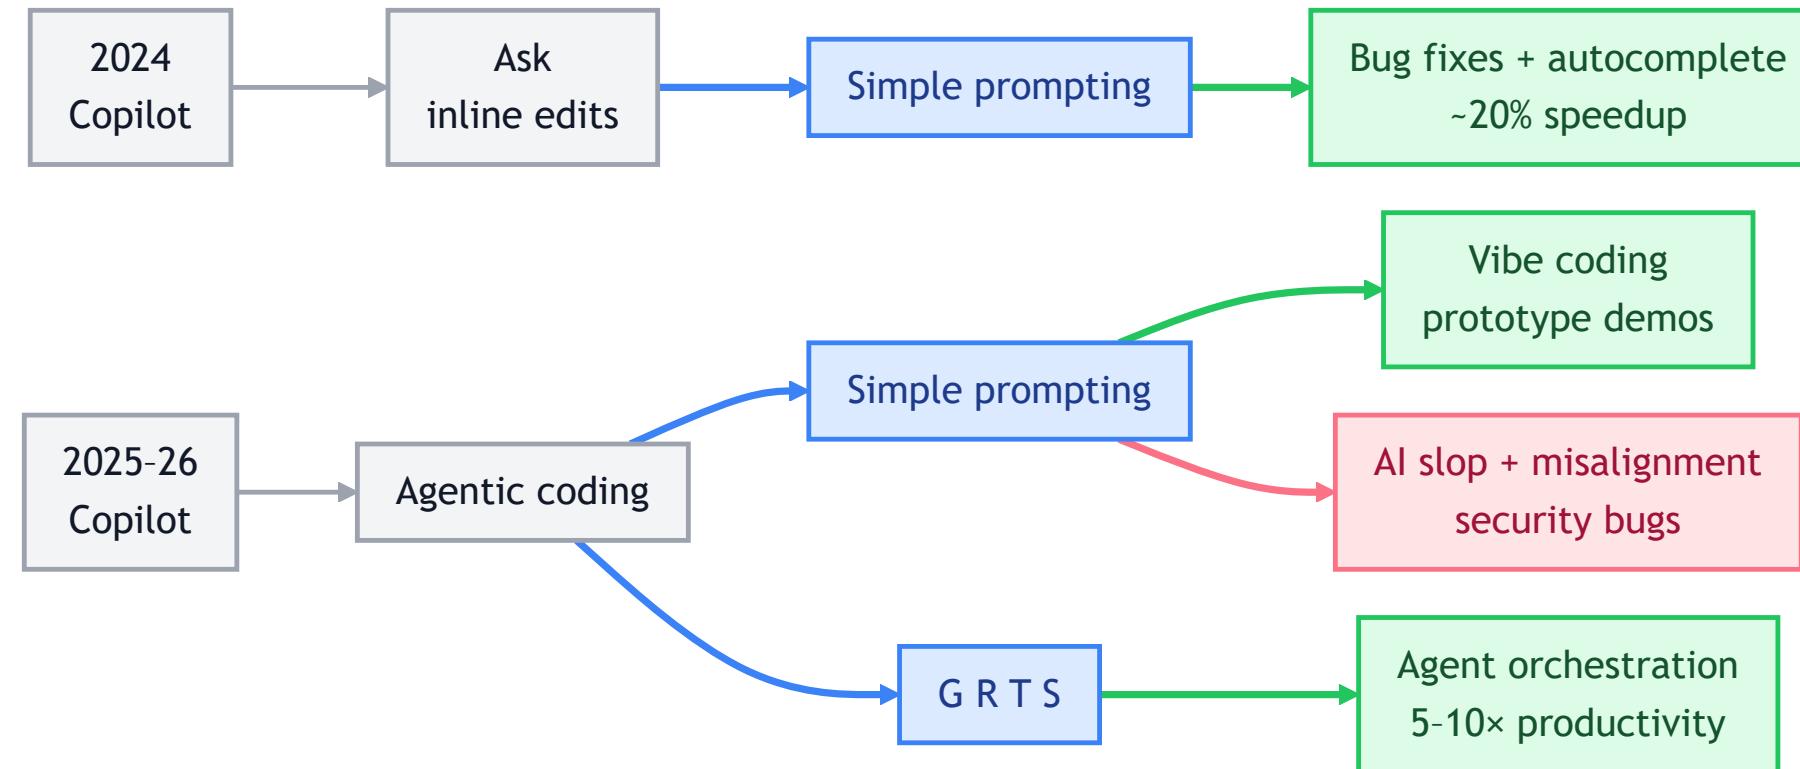




GR_TS

Don't build features. **Build the system that builds features.**

Coding Agent: The very eager stupid genius



You will not be replaced by an agent...
but you may be replaced by someone who can direct one well.



- Get familiar with the eccentricities of the stupid genius
 - **First steps (concrete, low context, low risk tasks):** Ask it to review your code, debug your software and research the next idea.
- If the agent falls down a Rabbit hole 🐰 ➔ ask:
 - Research the best 3 approaches to this issue with pros/cons — cite evidence.
 - Check out this site <stack overflow link>. Does it help?
- If the agent tries to cheat or take a shortcut ➔ ask:
 - What is the proper way to fix this?
 - Right place for the code? Clean architecture? Utilities library?
- Always define “good”:
 - **Verification vs Validation** — built the thing right vs build the right thing ← we need this!
 - If possible, use an oracle / golden sample / explicit success criteria.

Takeaway: Spend 1-2 weeks getting a feel for working with the latest LLM's.



- TODO - keep reviewing from here forward!
- As of Dec 2025, the new models are able to hold 8 things in their mind at the same time - they are limited by the context you give them, not their intelligence.
- Build repo overview + working rules (.github/copilot-instructions.md) ← **Copilot can help with this!**
- Cite docs: build scripts, library locations, standards, linting guides
- For the top 3–8 areas, define “roles” with repo-specific best practices (custom agents + optional .github/instructions/*.instructions.md)
- Define boundaries: excluded (secrets), auto-approved (git fetch/status), and restricted files (CI/CD, security)

Tips:

- Keep instruction files short (100–300 lines). Refine or split as needed.
- Spend 1–2 days iterating; this pays off.
- When prompts miss something (edge cases, wrong file, global var), explain the failure and update the instruction file



Tools Capabilities | Abilities | Touch the external world



Without extra tools, agents are limited to reading and writing files in your repo.

- Issue read and write: [GitHub MCP](#) [Atlassian MCP](#)
- Accurate API and function calls: [Context7](#) + research online
- Semantic repo navigation on large codebases: [Serena](#)
- Deterministic actions: generate a script for repetitive/complex work (e.g., convert NUnit3 tasks to NUnit4 across many files)
- Truth sources: diffs, test output, CI checks, benchmarks/sim logs
- VS Code tasks: Build, test, lint, format, etc.
- Shell commands: git status, list files, get file content
- Tip: Use settings.json ->"chat.tools.terminal.autoApprove" for read-only commands



Specifications

Objectives | Checklists | Planning | Requirements



- You mostly review outputs and spot-check code/decisions.
- For larger work, spend time planning to reduce churn and wrong turns.
- **Level 1:** For bugs and scoped changes, we can just make a single prompt
- **Level 2:** 1–3 planning prompts (research/explain/3 options) → then implement
- **Level 3:** Create a checklist in Markdown → review → implement
- 1-page templates help a lot: [Task spec](#) + [rules of engagement](#)
- **Level 4:** Use [OpenSpec](#) or [Spec Kit](#): specify → clarify → plan → analyze → implement

A1: Multi-agent workflows

- You'll often be waiting on agents. Use that time:
 - Research the next feature in parallel
 - Use git worktrees / background agents for parallel branches
 - Delegate to cloud agents (ensure CI/CD is set up)
 - Split roles: Planner → Implementer → Reviewer

A2: Risk -> Gating | Ambiguity -> Planning

The good	The bad and the ugly
<p>reason about the goal plan steps take actions (read/edit/run tools) iterate until done (or stuck)</p>	<p>confidently wrong misses hidden constraints misapplies “best practices” thrashes without a stable hypothesis misuses tools (wrong env/partial runs)</p>

	Risk: Low	Risk: High
Ambiguity: Low	Great for agents (docs, refactors, small tests)	Needs gates + review (small but critical changes)
Ambiguity: High	Clarify first (spec + oracle)	Human-first (architecture/safety-critical/unclear bugs)

A3: References (Core)

- **G** — [Building effective agents \(Anthropic\)](#) — Designing agent loops: checkpoints, tool feedback, stop conditions.
- **R** — [Security \(VS Code Copilot\)](#) — Trust boundaries, tool approvals, prompt injection risks.
- **R** — [Workspace Trust \(VS Code\)](#) — Restricted Mode and why it matters for agents.
- **R** — [LLM01:2025 Prompt Injection](#) — Threat model + mitigations.
- **R** — [About GitHub Copilot coding agent](#) — Capabilities, limits, and governance.
- **T** — [Tutorial: Work with agents in VS Code](#) — Local/plan/background/cloud agent workflows + worktrees.
- **T** — [Use tools in chat \(VS Code\)](#) — [Tool approval](#) — Tool approvals, URL post-approval, and auto-approval tradeoffs.
- **S** — [Spec Kit](#) — Spec → Plan → Tasks to make “done” measurable.
- **S** — [CI \(GitHub Actions\)](#) — CI as a repeatable verification oracle.
- **S** — [A Minimal, Reproducible Example](#) — Make bugs/tasks reproducible.
- **S** — [Responsible use of GitHub Copilot coding agent](#) — Scope, acceptance criteria, review gates.

A4: References (More)

- G — Lessons from Anthropic (secondary write-up)
- T — GitHub Copilot in VS Code
- T — Get started with GitHub Copilot in VS Code
- T — Asking GitHub Copilot questions in your IDE
- T — Review AI-generated code edits (VS Code)
- R — Adding repository custom instructions for GitHub Copilot
- T — Context7 (GitHub)
- T — Serena docs
- T — git-worktree documentation
- S — About issue and pull request templates (GitHub)
- S — Configuring issue templates for your repository (GitHub)
- T — GitHub Copilot Workspace (GitHub Blog)
- T — What is Foundry Agent Service?